

Microkernels, Genode and Gapfruit OS

Timo Nicolai and Alice Domagala

2023

Gapfruit AG

- Microkernels
 - What is a Microkernel?
 - Advantages of Microkernels
 - Examples of Microkernel Systems
 - (L4) Microkernel History
- Genode OS Framework
 - What is Genode?
 - Capabilities
 - The Component Hierarchy
 - The C++ Runtime
 - A Simple Genode Component
- Gapfruit OS



Alice Damage

- Information Technologies Expert @ Epitech Paris
- Operating systems engineer @ Gapfruit
 - Framework updates and upstream contributions
 - Involved in platform support for various hardwares and architectures including x86_64 and arm_v8a
 - Involved in creation and maintenance of native, and Linux ported, device drivers



Timo Nicolai

- M.Sc. Information System Technology @ TU Dresden
- Previously intern @ Kernkonzept
 - Worked on Fiasco.OC/L4Re
 - Extended uvmm hypervisor
- Now operating systems engineer @ Gapfruit
 - Integrating Gapfruit OS with Azure IoT
 - Involved in designing and implementing our PKI



What is a Microkernel?

- Kernel only implements minimal functionality
- Drivers, resource management etc. implemented in userspace
- Kernel provides *mechanisms* instead of *policy*



Advantages of Microkernels

- **Trustworthiness**
 - Small trusted code base
 - Exploited drivers do not compromise kernel
- **Reliability**
 - Misbehaving components can be restarted without rebooting
- **Resource Management**
 - No “overpromising” of resources
 - Resource trading



Examples of Microkernel Systems

- **MINIX** [1]
- **GNU Hurd** [2]
- **Zircon/Fuchsia** [3]
- **Redox** [4]
- **L4 Family**



(L4) Microkernel History: Predecessors

- **RC 4000 Multiprogramming System** [5] (Per Brinch Hansen, 1969)
 - Possibly the first microkernel system
- **Mach** [6] (Richard Rashid and Avie Tevanian, 1985)
 - UNIX compatible
 - Basis of GNU Hurd
 - Painfully slow IPC
- **L4** [7] (Jochen Liedtke, early '90s):
 - Handwritten in Assembly
 - Fast IPC: synchronous, uses registers, cache friendly

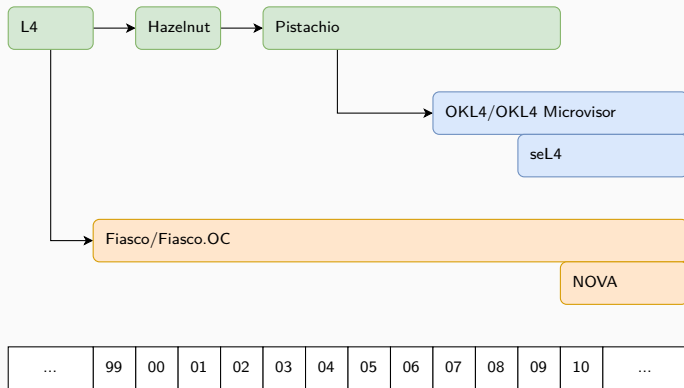


(L4) Microkernel History: Recent Developments

- **University of Karlsruhe** (Jochen Liedtke)
 - *L4Ka::Hazelnut* [8]: Written in C++
 - *L4Ka::Pistachio* [9]: Cross-platform
- **UNSW/NICTA** (Gernot Heiser)
 - *OKL4, OKL4 Microvisor* [10] (now @ *General Dynamics*)
 - *seL4* [11]: Formally verified, fast and suitable for realtime
- **TU Dresden** (Hermann Härtig)
 - *Fiasco/Fiasco.OC* [12] (now @ *Kernkonzept*): L4Linux
 - *NOVA* [13] (now @ *BedRock Systems*): Focus on virtualization



(L4) Microkernel History: Recent Developments



Based on [14]



(L4) Microkernel History: The Third Generation

- Capability-based security
- Virtualization
- Formal verification



Genode: What is Genode?

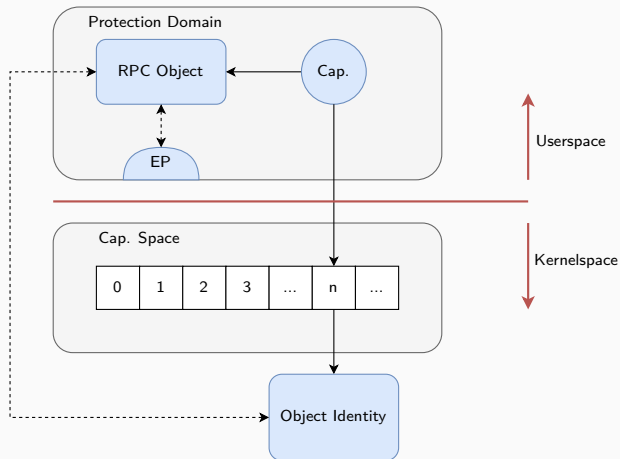
- Open source, developed at *Genode Labs* (+ *Sculpt* [15])
- Operating system *framework* [16]
- Includes the *HW* microkernel
- Runs on top of Linux, HW, Fiasco.OC, NOVA, OKL4, seL4



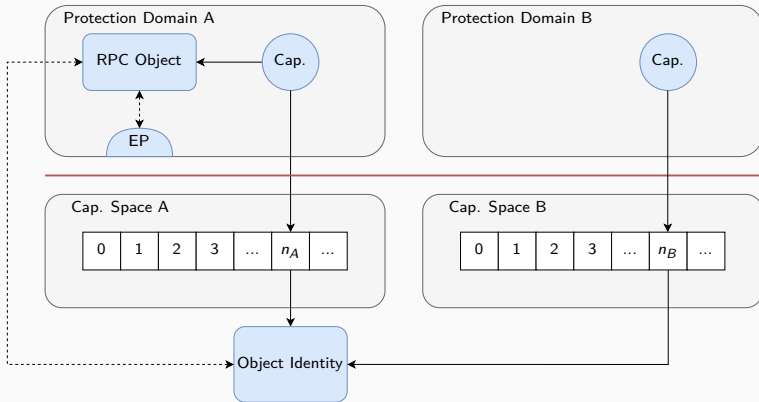
- Components need to provide/consume services
- Capabilities:
 - Live inside *protection domains*
 - Combine "pointers" and "access rights" to RPC objects
 - Are used to invoke methods of RPC objects



Genode: Capability Creation



Genode: Capability Creation



Genode: Capability Invocation

- Terminology: *client* (**B**) and *server* (**A**)
- RPC is mediated by kernel:
 1. Thread in **B** *calls* method of RPC object using its capability, providing client-local name, opcode, and arguments
 2. Kernel checks validity of capability
 3. Kernel finds corresponding RPC object in **A** and wakes up entrypoint thread
 4. Entrypoint thread in **A** finds and invokes the method of the RPC object
 5. Entrypoint thread in **A** may produce return value and *reply*
 6. Return value is passed back, if necessary, to calling thread which is woken up



Genode: Capability Delegation

- Need to be able to *delegate* capabilities between PDs
- Capabilities can be passed as RPC arguments/return values
- Kernel will create capabilities in receiving PD as necessary
- So, in order to transmit capabilities we require capabilities???



Genode: The Component Hierarchy: Parents and Children

- Every component except *core* has exactly one parent
- Parent provides children with resources and controls them
- Children start out with a single capability to parent
- Children have to trust parent but not the other way around

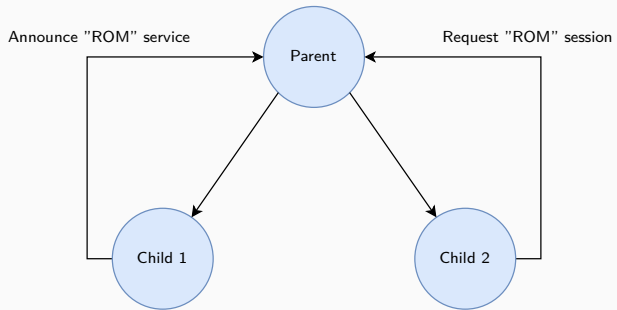


Genode: The Component Hierarchy: Services and Sessions

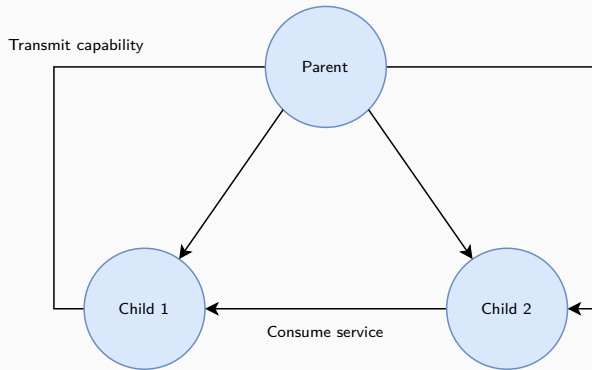
- Children announce *services* via RPC objects implementing *root interface*
- Other children can perform *session requests*
- Parent can modify these and deny them, handle them, forward them to a child/its own parent
- Once a session is established, communication is direct
- Clients need to trust servers but not the other way around
- Clients can share resources when initiating a session



Genode: The Component Hierarchy: Services and Sessions



Genode: The Component Hierarchy: Services and Sessions



- Most native Genode components don't use the STL
- Genode supports exceptions
- Genode comes with its own minimal C++ runtime



Genode: A Simple Genode Component

Let's try to build this:

```
timer tick 1
```

```
timer tick 2
```

```
timer tick 3
```

```
timer tick 4
```

```
timer tick 5
```



Genode: A simple Genode component

```
/* component which outputs a message 5 times and then exits */
class Main {
private:
    Genode::Env &_env;

    /* timer */
    Timer::Connection _timer { _env };
    /* timer tick counter */
    Genode::size_t _timer_ticks { 0 };
    /* timer trigger callback */
    Genode::Signal_handler<Main> _timer_handler { _env.ep(), *this, &Main::_handle_timer };

    void _handle_timer() {
        if (++_timer_ticks > 5)
            _env.parent().exit(0);

        Genode::log("timer tick ", _timer_ticks);
    }

public:
    explicit Main(Genode::Env &env) : _env { env } {
        /* initialize timer trigger handler */
        _timer.sigh(_timer_handler);
        /* trigger timer periodically (once every second) */
        _timer.trigger_periodic(1'000'000);
    }
};

/* component entry point */
void Component::construct(Genode::Env &env) {
    static Main main { env };
}
```



- Trustworthy/reliable OS for IoT edge gateways
- Collect data from devices and bring it to the cloud
- Control devices via the cloud



- Focus on integrity, stability, reliability and uptime
- Minimize human interaction with the system (*Zero-Touch*)
- Deploy devices and establish cloud connectivity automatically



- > 99% reduction of attack surface
- Strong control over service dependency and delegation
- Able to recover malfunctioning components (via *heartbeats*)
- Updating e.g. drivers without downtime is trivial
- Use TPM-backed PKI for automatic mass deployment



Gapfruit OS: Azure Example

Example *Device Twin*:

```
"properties": {  
  "desired": {  
    "dynamic": {  
      "deploy": [  
        {  
          "start": {  
            "name": "http_server",  
            "pkg": "gapstage/pkg/http_server/2023-01-01"  
          }  
        },  
        {  
          "start": {  
            "name": "mqtt_bridge",  
            "pkg": "gapstage/pkg/mqtt_bridge/2023-02-01"  
          }  
        }  
      ]  
    }  
  }  
}
```



Thank you for listening, are there any questions?



References

- [1] *MINIX 3*. <https://www.minix3.org/>.
- [2] *GNU Hurd*. <https://www.gnu.org/software/hurd/>.
- [3] *Fuchsia*. <https://fuchsia.dev/>.
- [4] *Redox OS*. <https://www.redox-os.org/>.
- [5] Per Brinch Hansen. *RC 4000 Software: Multiprogramming System*. Regnecentralen. 1969.
- [6] Richard Rashid, Avadis Tevanian, and Michael Young. "Mach: A New Kernel Foundation for UNIX Development". In: *USENIX Summer Conference*. 1986.
- [7] Jochen Liedtke. "Improving IPC by Kernel Design". In: *SIGOPS Oper. Syst. Rev.* 1993.
- [8] *L4Ka::HazelNut*. <https://www.l4ka.org/57.php>.
- [9] *L4Ka::Pistachio*. <https://www.l4ka.org/65.php>.
- [10] Gernot Heiser and Ben Leslie. "The OKL4 Microvisor: Convergence Point of Microkernels and Hypervisors". In: *Proceedings of the First ACM Asia-Pacific Workshop on Workshop on Systems*. 2010.
- [11] Gerwin Klein et al. "SeL4: Formal Verification of an OS Kernel". In: *Proceedings of the ACM SIGOPS 22nd Symposium on Operating Systems Principles*. 2009.
- [12] *The L4Re Microkernel*. <https://os.inf.tu-dresden.de/fiasco/>.
- [13] *NOVA Microhypervisor*. <https://hypervisor.org/>.
- [14] Gernot Heiser and Kevin Elphinstone. "L4 Microkernels: The Lessons from 20 Years of Research and Deployment". In: *ACM Trans. Comput. Syst.* (2016).
- [15] *Sculpt OS*. <https://genode.org/download/sculpt>.
- [16] *Genode Foundations*. <https://genode.org/documentation/genode-foundations/index>.

